

Vector DB (Search) Technical Specification

v3.5

Decision-Oriented | Hardware-Aware | Evolutionary Narrative
Authors: Data Motifs, Gemini, ChatGPT

0. Notation Table

Symbol	Meaning
N	Number of stored vectors
D	Vector dimensionality
B	Bytes per element
n	Number of vectors touched (algorithmic skipping)
n_probe	Number of vectors probed in IVF cluster
D_eff	SIMD-aligned effective dimension
B_quant	Bytes per element after quantization
Γ	Locality multiplier (hardware)
γ_{NUMA}	NUMA factor (cross-socket memory)
Λ	Software / glue code overhead
Parallelism	Number of CPU cores / threads
Φ	Unified system strain
$\Phi_{\text{distributed}}$	Distributed system strain
InstrCost	CPU instruction cost per operation (function of quantization)
K	Number of top results requested
S	Number of shards / nodes
χ	Contention factor (<1) for concurrent queries

D_eff — SIMD-aligned effective dimension. Usually rounds up to nearest SIMD lane multiple.

Example: If $D = 128 \rightarrow D_{\text{eff}} = 128$; if $D = 100 \rightarrow D_{\text{eff}} = 112$ or 128 depending on SIMD width (AVX, AVX512, etc.).

1. Historical Foundation: Unit Cost of Search

- **Computational Cost:**

$$C_{Search} = N \cdot D \cdot MathCost$$

- **Memory Transfer Cost:**

$$M_{Total} = N \cdot D \cdot B + Metadata$$

Observation:

- Naive linear search (K-NN) is $O(N)$, memory-bound $\rightarrow$ Memory Wall.
 - All later equations refine these fundamental costs.
-

2. Evolutionary Derivation: Algorithmic Costs

2.1 Naive Linear Search

$$\begin{aligned} C_{Unit} &= D \cdot MathCost + D \cdot B \cdot MemoryAccessCost \\ C_{Total} &= N \cdot C_{Unit} \end{aligned}$$

- Linear in N , dominated by memory for large datasets.

2.2 Clustered Indexing (IVF)

$$C_{Total} = k \cdot C_{Unit} + n_{probe} \cdot \frac{N}{k} \cdot C_{Unit}$$

- k = number of clusters
- n_{probe} = number of clusters searched (tune for recall α)
- Reduces effective $N \rightarrow n$, adds centroid search overhead.

2.3 Graph-Based Search (HNSW)

$$C_{HNSW} \approx efSearch \cdot \log(N) \cdot C_{Unit}$$

- $efSearch$ = number of entry points explored (high recall $\alpha \rightarrow$ higher $efSearch$).
- Visits only relevant neighbors; logarithmic cost with constant factor dependent on $efSearch$.

2.4 Algorithmic Skipping & Clustering

- Skipping factor: $n \ll N$
- Partition vectors into clusters (IVF / HNSW)
- Efficiency Ratio (ER):

$$ER = \frac{DataUsed}{DataMoved}$$

- Reduces search cost from $O(N) \rightarrow O(\sqrt{N})$ or $O(\log N)$.

2.5 Representation Evolution: Quantization

Method	Bytes / Dim	Effect
SQ	1	4× memory savings, minor accuracy loss
PQ	0.02–0.1	50× memory savings, fits billion-scale datasets
BQ	0.125	Ultra-fast XOR math; high throughput (512-bit / cycle)

- SQ: reduces B
- PQ: splits $D \rightarrow D_{sub}$, reduces memory footprint
- BQ: compresses both B & D, speeds computation
- Instruction cost varies with quantization:

$$InstrCost = f(B_{quant})$$

3. Hardware & Software Awareness

3.1 Locality Multiplier (Γ)

Access Pattern	Γ Range	Notes
Sequential / Flat Scan	1.0	Prefetching effective
Clustered (IVF)	2–5	Sequential within clusters; some cache misses
Random / Graph (HNSW)	5–10	Pointer-chasing; CPU stalls
NUMA-crossing	+1–2×	γ_{NUMA} : cross-socket memory increases effective Γ

3.2 Glue Code & Software Overhead (Λ)

Source	Latency / vector	Notes
Serialization	1–10 μs	Memory mapping helps
Context Switching	5–50 μs	Pin threads to reduce OS hops
Network / RPC	50–500 μs	Cluster-dependent; consider percentile latency

Observation:

- Λ dominates small datasets ($N < 10k$).
 - Total overhead scales linearly with vectors touched.
-

3.3 Unified System Strain (Φ)

$$\Phi = \Lambda + \frac{\Gamma \cdot \gamma_{NUMA} \cdot (n \cdot D_{eff} \cdot B_{quant} \cdot InstrCost(B_{quant}))}{Parallelism \cdot \chi}$$

Where:

- n = vectors touched (skipping applied)
- D_{eff} = SIMD-aligned dimensions
- B_{quant} = bytes per element after quantization
- $InstrCost(B_{quant})$ = instruction cost per operation, depends on quantization
- $Parallelism$ = CPU cores / threads
- γ_{NUMA} = NUMA locality factor
- $\chi < 1$ = concurrency degradation

Interpretation:

- $\Lambda \rightarrow$ software/network latency
- $\Gamma \cdot \gamma_{NUMA} \rightarrow$ hardware memory penalty
- $\chi \rightarrow$ throughput reduction under concurrent queries
- $Fraction \rightarrow$ parallelized vector computation
- $\Phi \rightarrow$ practical metric for real-world deployment

InstrCost — CPU instruction cost per operation (Bitwise > Integer > Float).

For **Binary Quantization (BQ)**, $InstrCost \approx 0.02 \times \text{float32}$, because XOR + POPCNT can process **512 bits per cycle** on modern CPUs.

This explicitly separates **compute throughput gains** from memory savings.

4. Distributed / Multi-Node Systems

$$\Phi_{Distributed} = \frac{\Phi_{Local}}{S} + (\Lambda_{Network} + Cost_{Merge}(K, S))$$

- Merge cost: $O(K \log S)$ (Top-K aggregation)
- $\Lambda_{Network}$ = network latency

- Multi-node tax arises once $N \gtrsim 100M$

Scaling Simulation (100M vectors):

Nodes (S)	RAM / Node	Latency Local	Latency Merge	Total
1	400 GB	500 ms	0 ms	500 ms
10	40 GB	50 ms	5 ms	55 ms
50	8 GB	10 ms	15 ms	25 ms

- Adding nodes beyond threshold increases merge/network latency.
- Profile for γ_{NUMA} and χ under concurrent load.

5. Practical Decision Table

Scale (N)	Bottleneck	Algorithm / Index	B / D	Γ	Λ	Notes
<10k	Software	Flat / In-Process	float32	1.0	Dominant	Flat scan faster; minimize network
10k–1M	Cache	IVF / Flat-GPU	int8	2–5	Moderate	Sequential cluster access improves Γ
1M–10M	Search n	HNSW	PQ	5–10	Minor	Algorithmic skipping critical; tune efSearch
10M–100M	RAM / Network	Disk-based IVF + Sharding	BQ	5–10	Significant	Monitor merge/network; γ_{NUMA} may spike
>100M	RAM & Network	Sharded HNSW + PQ/BQ	PQ/BQ	5–10	High	Plan $\Phi_{\text{distributed}}$; consider χ ; hierarchical merge & caching

6. Verification Checklist

- SIMD alignment: D_{eff} aligned to 16/32 (AVX2/AVX512)
- Memory bandwidth: $\text{Data size} / \text{latency} \leq \text{RAM bandwidth}$
- Cache fit: Index vs L3 cache
- Glue check: Λ not dominating small N
- Quantization: B_{quant} matches memory/accuracy trade-offs
- Parallelism & contention: cores/threads sufficient; tune χ
- Multi-node: validate $\Phi_{\text{distributed}}$, merge/network, NUMA effects

7. Cost Modeling

$$Monthly_Cost = N \cdot D_{eff} \cdot B_{quant} \cdot GB_Rate$$

- Memory dominates $\rightarrow$ prioritize reducing B_{quant}
-

8. Decision Flow

1. Input: N , D , B , Target Latency
 2. Compute D_{eff} , size
 3. Compute $\Phi_{min} = \text{Size} / \text{Bandwidth}$
 4. Check $N < 15k$? $\rightarrow$ Flat / In-Process
 5. Select index based on α (Recall)
 6. Adjust n , B_{quant} , Γ , Parallelism, χ
 7. Compute Φ and verify $\leq$ Target
 8. Deploy Locally or Distributed ($\Phi_{distributed}$)
-

9. Evolutionary Summary

- Unit Cost $\rightarrow$ Memory Wall
 - Algorithmic Skipping $\rightarrow$ reduce effective N
 - Quantization $\rightarrow$ reduce $D \times B$
 - Hardware Awareness $\rightarrow \Gamma$ & γ_{NUMA}
 - Software Awareness $\rightarrow \Lambda$
 - Unified System Strain $\rightarrow \Phi$ combines all variables
 - Cluster / Shard Extension $\rightarrow \Phi_{distributed}$
-

10. Notes / Recommendations

- Always simulate Φ before deployment (choose shard count, quantization levels)
 - Datasets $> 10M \rightarrow$ PQ or BQ essential
 - Control Λ for small N ; avoid unnecessary network layers
 - Hardware tuning (SIMD, cache, RAM bandwidth) directly impacts Φ
 - Document N , D , B , α , resulting Φ
-

11. Minor Caveats / Considerations

- Γ ranges for HNSW (5–10) are empirical; profile for your CPU & RAM
 - Λ dominance less significant for fully optimized C++ or GPU in-process searches
 - Distributed merge $O(K \log S)$ is naive; pre-aggregation can reduce cost
 - Parameter tuning (n_{probe} , k , efSearch , quantization) critical for Φ
 - Edge cases: small D vectors, GPU / mixed setups not modeled in Φ
 - Network variability & memory bandwidth saturation can cause nonlinear effects
-

12. Decision Flow (Sanity Check)

Add a **bandwidth saturation check** to ensure realistic Φ predictions under high query rates:

$$\text{If } (n \cdot D_{\text{eff}} \cdot B_{\text{quant}} \cdot \text{QPS}) > \text{Total RAM Bandwidth} \Rightarrow \chi \downarrow$$

- **QPS** = Queries per second
- χ = Contention factor, reduces effective parallelism when memory bandwidth is saturated
- Ensures multi-query contention is captured before deployment

Decision Flow:

1. Input: N , D , B , Target Latency, QPS
 2. Compute D_{eff} , Size
 3. Compute $\Phi_{\text{min}} = \text{Size} / \text{Bandwidth}$
 4. $N < 15k$? $\rightarrow$ Flat / In-Process
 5. Pick index & quantization (target α / recall)
 6. Adjust n , B_{quant} , Γ , Parallelism, χ
 7. **Sanity check:** bandwidth vs QPS $\rightarrow$ adjust χ
 8. Compute Φ (or $\Phi_{\text{distributed}}$)
 9. Deploy locally or distributed
-

13. Visualizing Trade-offs (Recommended for Stakeholders)

- Create “**Recall vs Latency**” frontier plots:
 - X-axis: Latency (ms)
 - Y-axis: Recall (α)
 - Lines/curves for Flat, IVF, HNSW, PQ, BQ
- Demonstrates why we move from flat scanning $\rightarrow$ clustered $\rightarrow$ graph-based $\rightarrow$ quantized
- Makes **trade-offs between accuracy, memory footprint, and compute visible** to non-technical stakeholders